

Before you let an agent drive a screen

01

Put it on its own machine

A desktop agent inherits every session already open on that desktop. Your mail, your VPN, your password manager. Give it a machine holding nothing you would not hand over.

02

Decide what a screenshot is

Computer use works by looking at the screen, and every screenshot is a copy of whatever was visible, sent off the machine. Anything that must not leave should never be rendered.

03

Name the applications it may open

An allowlist of programs is a harder boundary than an instruction not to open things. An instruction is advisory. A missing binary is not.

04

Put the irreversible clicks behind a person

A screen has no permission model, only pixels. Send, delete, submit and pay all look exactly like clicking, so the boundary has to be yours rather than the interface's.

05

Cap the session, not the outcome

A stuck computer-use agent looks identical to a working one, because it is still clicking. Limit the minutes and the number of actions, since there is no error to catch.

06

Assume the page will change

A UI agent is coupled to a layout nobody promised to keep. Decide now what it should do when the button moves, because the button will move.

**Save this before someone
on your team tries it.**

- Ali Lari · AI systems for founders