

Before you let an agent drive a screen

01

Put it on its own machine

It inherits every session already open on that desktop: your mail, your VPN, whatever your password manager has unlocked. Give it a machine with nothing on it you'd mind handing over.

02

Decide what a screenshot is

It works by looking at the screen, so every screenshot is a copy of whatever was visible, sent off the machine. If something can't leave the building, don't render it.

03

Name the apps it's allowed to open

An allowlist of programs holds better than an instruction not to open things. You can talk it out of an instruction. You can't talk it into a binary that isn't installed.

04

Put the irreversible clicks behind a person

A screen has no permission model, just pixels. Send, delete, submit and pay all look exactly like clicking, so the boundary has to come from you rather than from the interface.

05

Cap the session, not the outcome

A stuck agent looks exactly like a working one, because it's still clicking. There's no error to catch, so cap the minutes and the number of actions instead.

06

Assume the page will change

You've coupled it to a layout nobody promised to keep. Decide now what it does when the button moves, because the button will move.

**Save this before someone
on your team tries it.**

- Ali Lari · AI systems for founders