
Copy this

How Meta built an agent you'd let near your own accounts

01

Give it its own box

Muse runs on a dedicated VM in the cloud, and no other agent can reach it. Your agent gets its own machine or container, not a folder on a shared server.

02

Put a second agent in charge of no

A separate Sentinel agent sits on the same machine, kept apart at the system level. Nothing reaches the internet unless it approves. Build the approver first, before the agent has anything to send.

03

Credentials it can use but never read

Muse has no visibility into passwords or payment methods. They sit in secure storage it can call. If your agent can print a password, it can leak one.

04

Ask before send, ask before pay

Muse checks with the person before an email goes out or money moves. Two actions, one rule: nothing irreversible without a yes.

05

Show what it did and what it plans to do

A complete audit trail, including the steps it hasn't taken yet. 'Completed successfully' is not a log. The before, the after and the plan are.

06

Scope per app, and let people cut it

For email, a person picks read-only or read-and-send. Any connection can be dropped at any time. Give every integration that same switch.

**Save this before you
connect an agent to your
email.**

- Ali Lari · AI systems for founders