

The NSA's distillation checks, translated for a product with one model behind it

01

Track the seat, not the login

Shared premium seats were the entry point in advisory AA26-251A. One account used by a whole team reads as one heavy user. Count sessions and devices per seat, and compare usage per seat to the median.

02

Flag day-one maximum usage

The agencies' own tell: a new account running at its ceiling from the first hour. Real users ramp. Alert on any new account above the 95th percentile in its first week.

03

Watch throughput, then topics

Enterprise-scale volume on a consumer plan is the first signal. Thousands of requests on one narrow subject is the second. Both are counts already sitting in your logs.

04

Log the model call, not just the request

MITRE ATLAS M0024, AI telemetry. Store prompt length, output length and which model answered. Without that you cannot see who is harvesting your reasoning traces.

05

Degrade before you block

For high-confidence cases the advisory recommends a downgraded model or subtle variation, because a hard block just triggers failover to the next account. Serve the cheaper model and keep watching.

06

Rate-limit per seat, then per pattern

MITRE ATLAS M0004. A cap per seat stops the casual case. A cap across seats that share registration details, an IP range or a schedule stops the organized one.

Save this before your usage report surprises you.

- Ali Lari · AI systems for founders