Oct 2026

AI agents are everywhere. The returns are not.

Six studies, one pattern, five moves for your business

THE SIX SOURCES

McKinsey | Wharton + GBK | OpenAI | Asana | Pax8 | Nylas

Swipe for the exhibits

Ali Lari

1 · Adoption

Large companies are scaling agents. Small ones have not moved

Share of respondents whose organization is scaling AI agents, McKinsey global survey, %



Large = more than $1 billion in annual revenue.

>50%

of output tokens in OpenAI's EU enterprise sample come from Codex and ChatGPT Work

8x

growth in weekly users of those tools at EU small and medium businesses since May

22–85%

agentic share of tokens, Slovenia to Lithuania

What it means

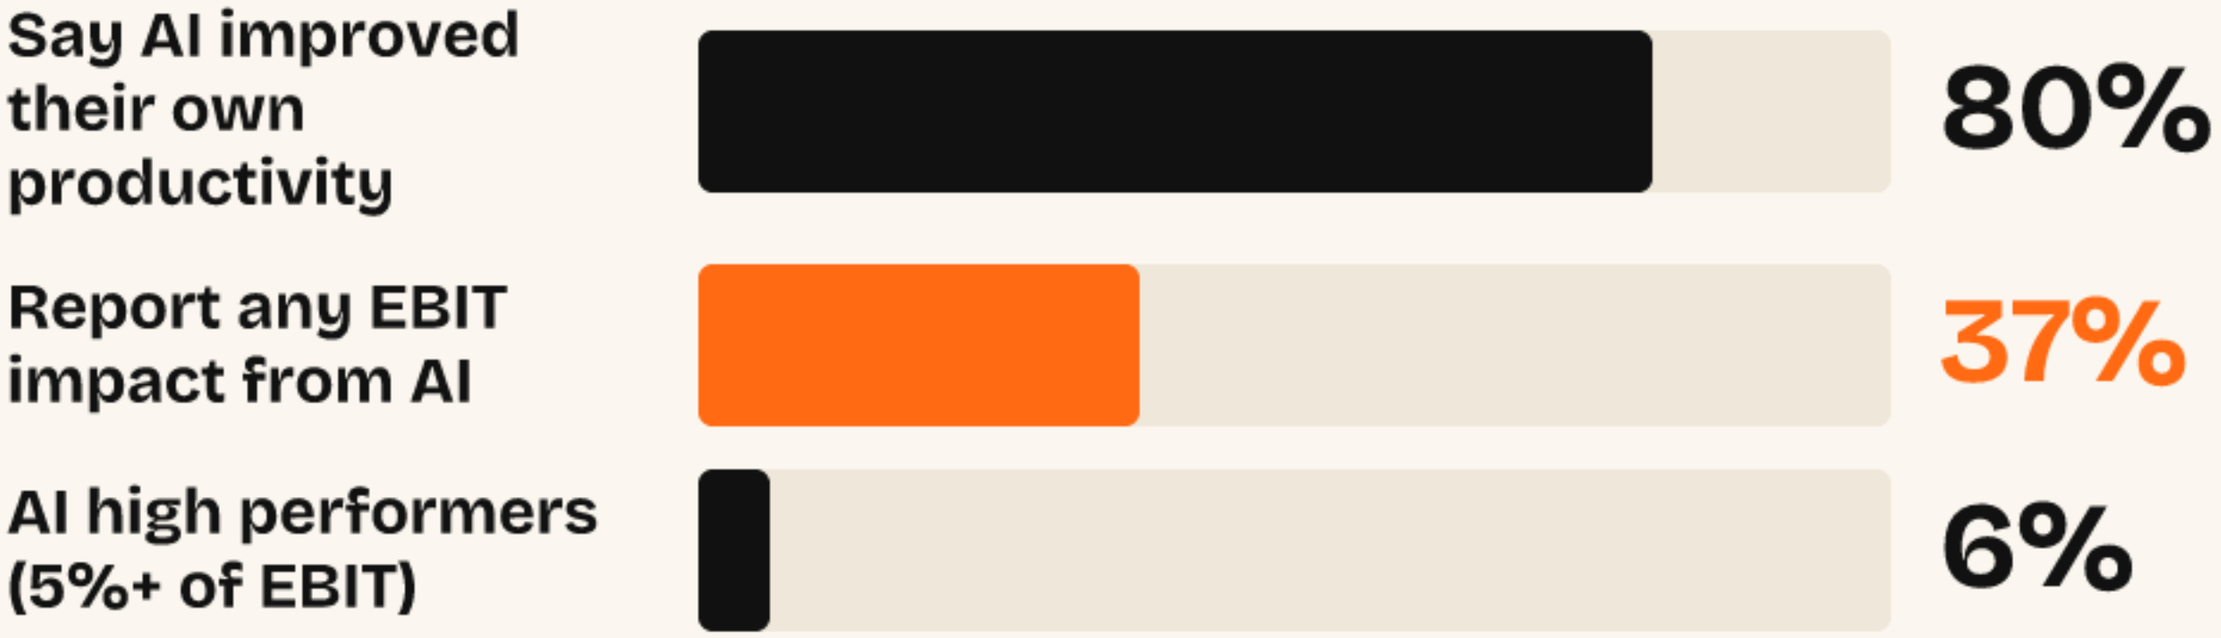
Agent use is concentrated in big firms, and OpenAI's figures describe its own customers. Read them as direction, not a market average.

Source: McKinsey, The state of AI in 2026, 25 Aug 2026 (respondent-reported). OpenAI Economic Research, The EU's next chapter in AI is agentic, 14 Oct 2026 (OpenAI's own ChatGPT Enterprise and Business accounts).

2 · Value

People say AI makes them faster. The P&L has not moved

Share of respondents, McKinsey global survey 2026, %



The EBIT share is unchanged from 2025.

60%

expect to increase AI investment next year

20%

limit AI use because of operating costs, tokens included

32%

chose not to buy software they could build with coding agents

What it means

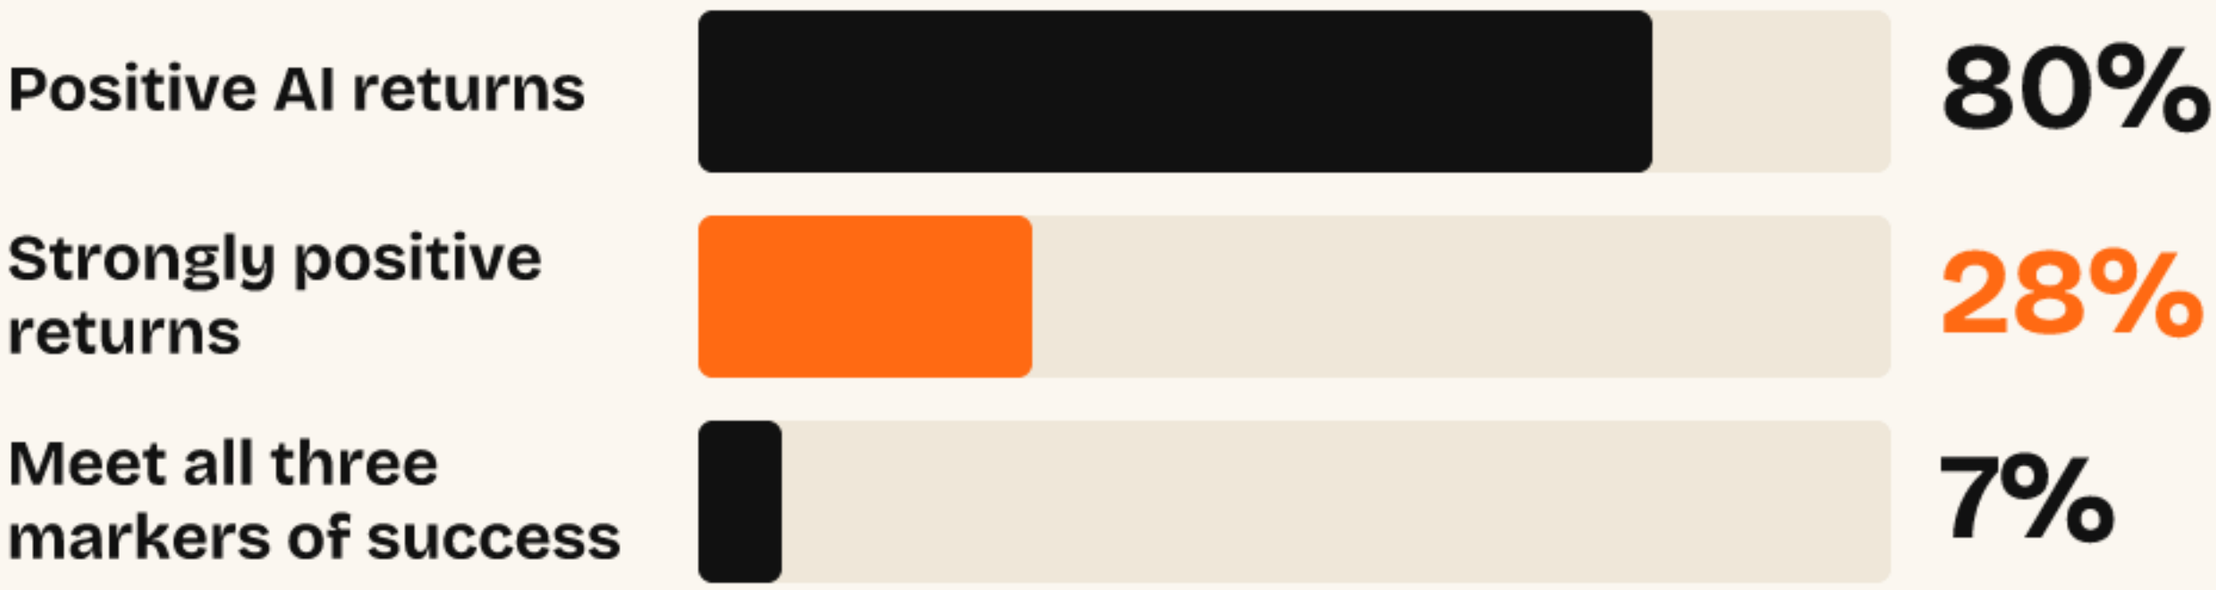
McKinsey says its high performers redesign workflows around AI instead of inserting it into existing ones.

Source: McKinsey, The state of AI in 2026: On the road to ROI, 25 Aug 2026. High performers: 5%+ of EBIT from AI and impact described as significant.

3 · Returns

Most leaders report positive returns. Few report strong ones

Share of enterprise leaders, Wharton and GBK, %



Markers: proving value, transforming workflows, aligning the organization.

40%

use hours saved as the ROI metric, the most common

27%

use net-new revenue, the least used major measure

What it means

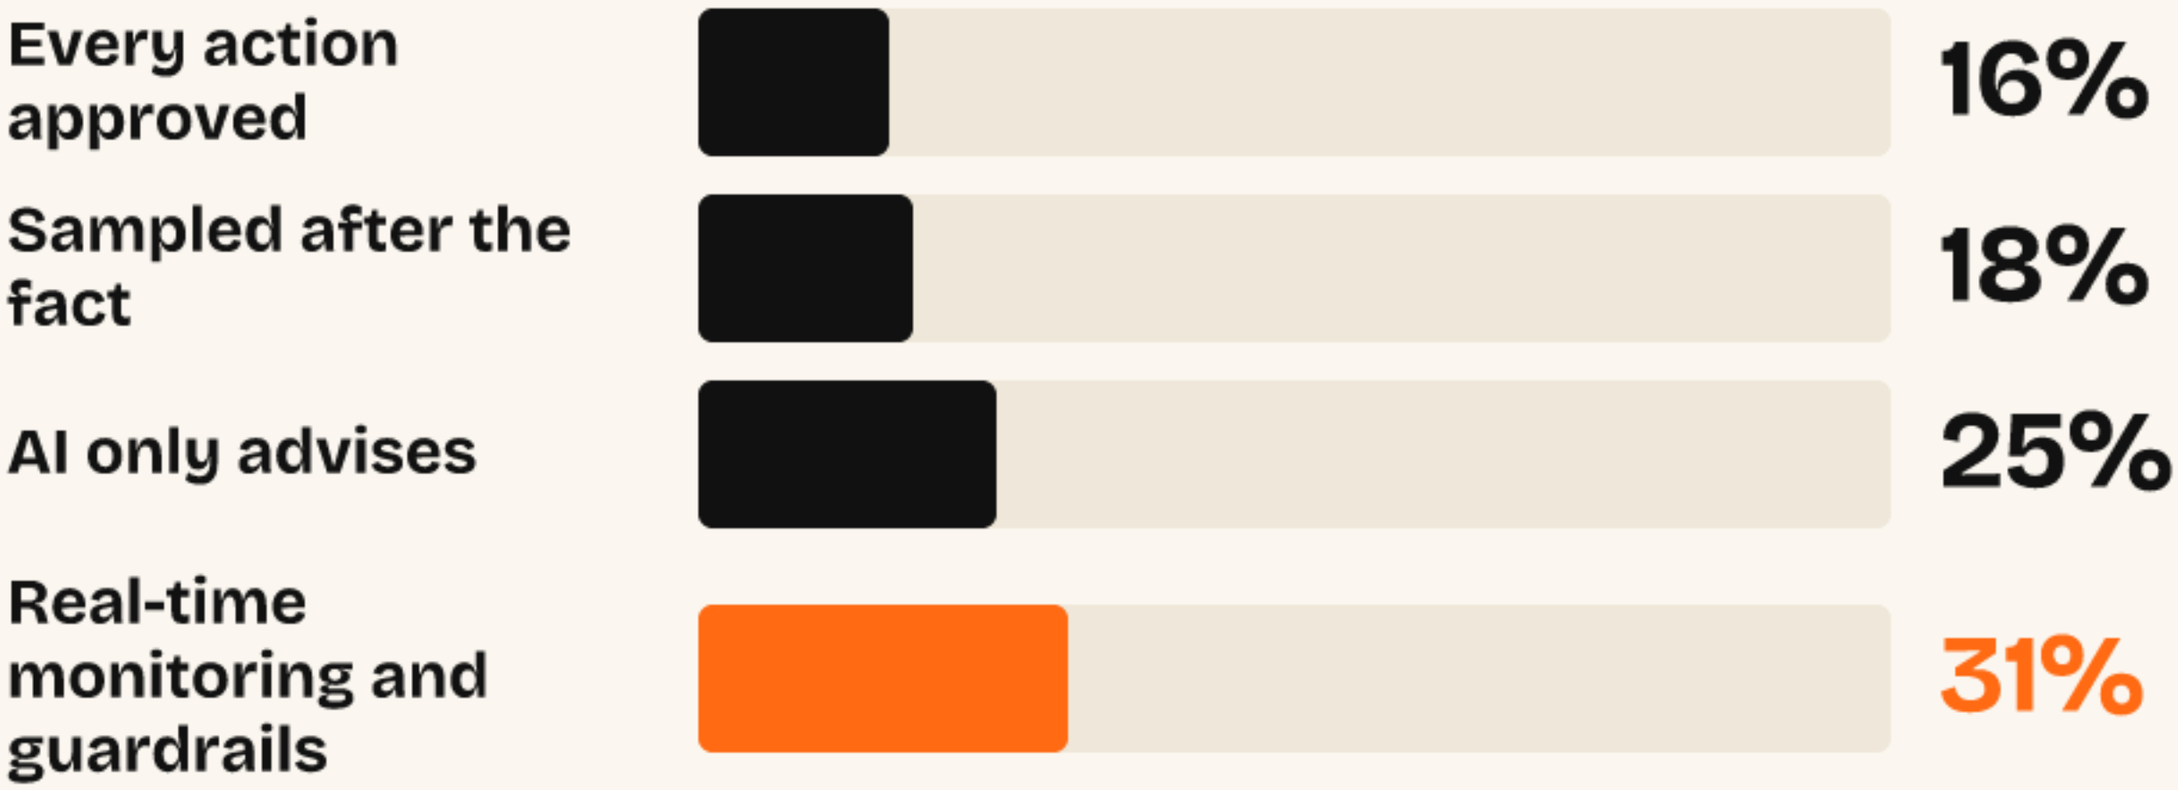
Departments tracking at least one growth measure report strong returns 35% of the time. Efficiency alone: 13%.

Source: Wharton Human-AI Research and GBK Collective, The Adaptive Advantage, 13 Oct 2026. U.S. enterprise leaders (n=1,007) and employees (n=409), fielded 9 to 29 Jul 2026.

4 · Oversight

Approving every action is where workflows change the least

Share of departments that have transformed their workflows, by oversight model, % (GBK)



The guardrails group is small (68 respondents): directional.

48%

of departments: AI does the work and a person approves before anything happens

17%

of departments let an agent finish a task without per-step approval

What it means

Paying for autonomy and then reviewing every action is the supervised middle. Monitoring with guardrails shows the most workflow change.

Source: GBK Collective write-up of The Adaptive Advantage and the Wharton report, 13 Oct 2026. U.S. enterprise respondents, fielded 9 to 29 Jul 2026.

5 · Reliability

The more agents a company runs, the more often one fails

Share of organizations reporting an agent incident in the past 90 days, by agents running, % (54% overall)



The 6 to 20 band is not broken out in the release.

31%

had an agent trigger a workflow the team had to reverse

25%

had one update the wrong record

18%

had one send an incorrect or inappropriate email

What it means

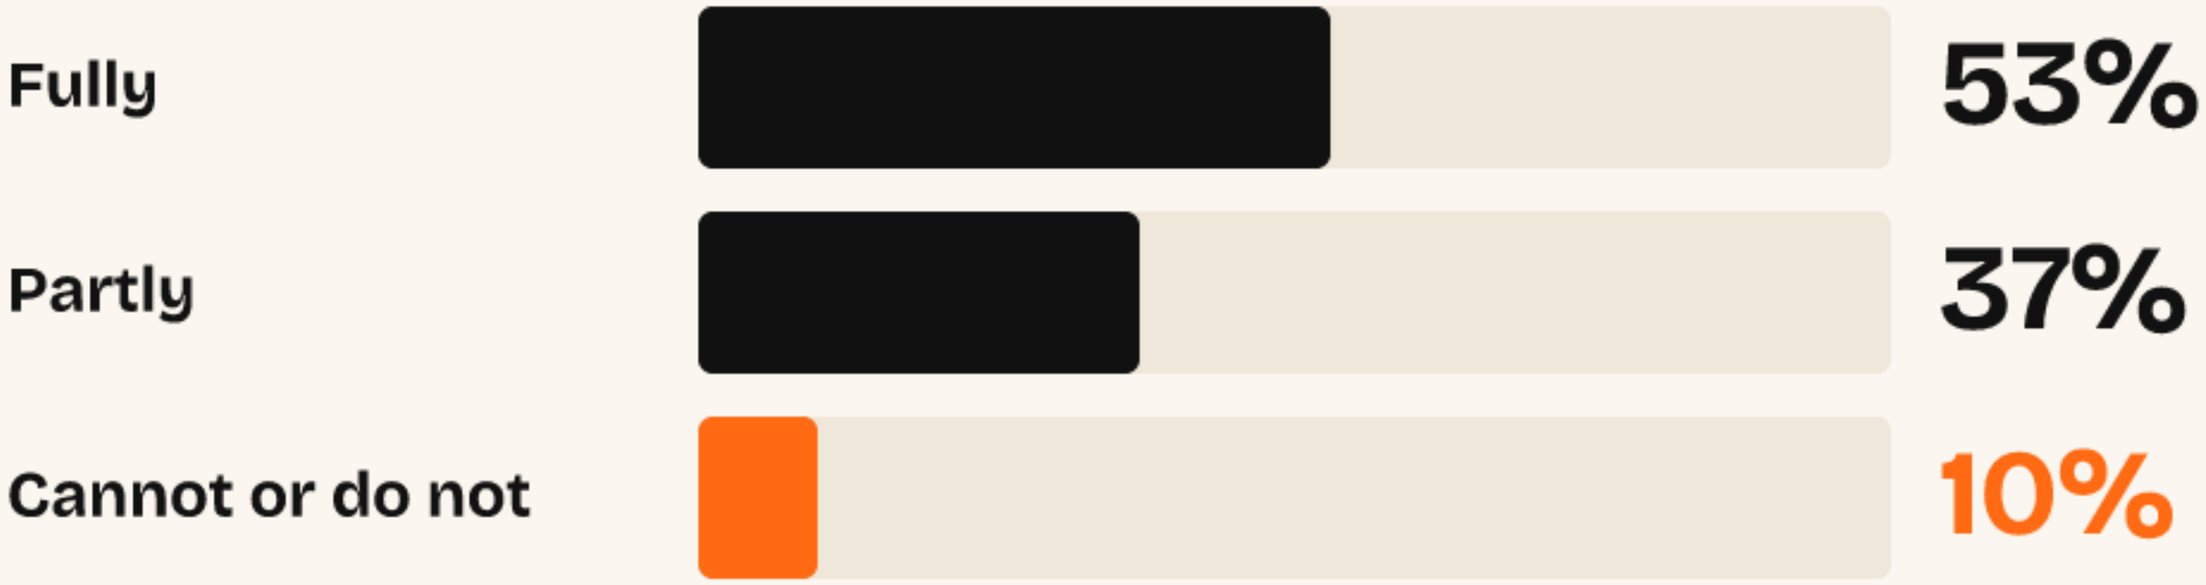
Incidents rise with the size of the fleet, and the most common one is a workflow a person has to undo.

Source: Nylas, The State of Agentic AI: H2 2026, 20 Oct 2026. U.S. professionals who build, buy or influence AI, n=1,030, fielded 26 Aug to 12 Sep. Vendor-run survey.

6 · Accountability

Only half of companies can fully audit their agents

Share who can audit what their agents did last week, %



26%

give agents a dedicated identity

66%

can fully audit where agents have their own identity

51%

can where agents use an employee's login

35%

of AI-engaged small firms found AI running without approval in six months

What it means

An agent that logs in as an employee leaves a log that names the employee. Identity decides whether an audit is possible.

Source: Nylas, 20 Oct 2026 (n=1,030, vendor-run; Nylas sells agent identity). Pax8 SMB AI Pulse Q3 2026, 12 Oct 2026 (n=1,150 SMB leaders in the U.S., U.K., Germany and Sweden).

7 · Coordination

Leaders count agents. Workers pay the coordination cost

Share using three or more agents in a typical week, Asana, %



23%

of workers say agents address the tasks they most need help with

46%

of workers using 2+ agents lose over 10% of their time to switching and reconciling

38%

say over a quarter of agent work is duplicated

What it means

Deployment is decided at the top and absorbed at the bottom. Asana's advice: measure finished work, not agent activity.

Source: Asana with Censuswide, AI Agents in the Flow of Work, 14 Oct 2026. 1,501 CIOs and AI leaders and 3,001 office workers in the U.S. and U.K., Sep 2026. Vendor-run; respondent estimates.

What to do with it

Five moves the numbers support

1. **Redesign one workflow before you add an agent**
   67% of leaders in departments meeting all three transformation conditions report strongly positive ROI, against 19% of the rest (Wharton, 13 Oct).

2. **Track a growth number next to hours saved**
   Strong returns: 35% with a growth measure, 13% with efficiency alone (Wharton).

3. **Swap sign-off on every step for monitoring and guardrails**
   Workflow change: 31% under real-time monitoring, 16% where every action is approved (GBK, directional).

4. **Give each agent its own login and a readable log**
   66% can fully audit agents with their own identity, 51% where they use an employee's login (Nylas, 20 Oct).

5. **Budget agent usage as a variable cost**
   49% of businesses that paid for AI got a bill above expectations (Pax8, 12 Oct); 20% limit use over cost (McKinsey).

Which of the five would you fix first?

Source: Slides 2 to 8 carry their sources. Pax8 SMB AI Pulse Q3 2026, 12 Oct 2026 (n=1,150 SMB leaders in the U.S., U.K., Germany and Sweden; vendor-run).